

# Project Report: Enchanted Emporium Gen-AI Assistant

## API Choice: Google Generative AI (Gemini)

We chose **Google's Generative AI (Gemini)** API for both:

- **Embeddings:** models/embedding-001 for semantic chunk representation.
- **Chat Model:** gemini-1.5-flash for fast and coherent LLM-based responses.

### Rationale:

- Seamless integration via langchain-google-genai.
- Supports both embeddings and chat models with the same API key.
- Gemini 1.5 models offer strong performance at lower latency (ideal for real-time assistant).

## Retrieval-Augmented Generation (RAG) Approach

### Pipeline:

1. **PDF Ingestion:** Extracted text using PyPDF2.
2. **Chunking:** Used RecursiveCharacterTextSplitter (chunk size: 1000, overlap: 200) to retain context.
3. **Vectorization:** Generated embeddings from each chunk using GoogleGenerativeAIEmbeddings.
4. **Storage:** Stored embeddings using **FAISS**, an in-memory vector database.
5. **Retrieval + QA:** On user query, the top-3 most relevant chunks are retrieved and passed as context to the Gemini chat model using RetrievalQA.

## Frontend Design (Streamlit)

- Built a user-friendly UI using **Streamlit**.
- Key features:
  - **PDF Upload Widget** for knowledge base ingestion
  - **Real-Time Chat Input** to ask about potions
  - **Dynamic Chat History Display** (last 5 queries)
  - **Sidebar Instructions** to guide first-time users

## Challenges Faced

1. **API Key Management:**
  - a. Ensuring the Gemini API key is securely passed and not hardcoded.
  - b. Alternative: Move to. stream lit/secrets.toml for production.
2. **Vector Index Memory Management:**
  - a. FAISS is fast but memory-bound; large PDFs may require indexing optimizations or persistent storage.
3. **PDF Quality:**
  - a. Extracting text from scanned or image-based PDFs failed silently (since `extract_text ()` returns `None`).
4. **Latency:**
  - a. While Gemini Flash is optimized, vector search + LLM inference still takes 1–3 seconds per query.
5. **Chunking Trade-offs:**
  - a. Small chunks increase relevance but risk context loss.
  - b. Larger chunks retain context but may dilute retrieval precision.

## Conclusion

The assistant successfully demonstrates a practical use of RAG with Gemini API to simulate an intelligent, domain-specific recommendation system. Its modular design and clean UI make it extensible for other domains (e.g., recipe books, tech manuals, FAQs)